

# Networks and Distributed Systems

## Lecture 15 – TCP sliding window

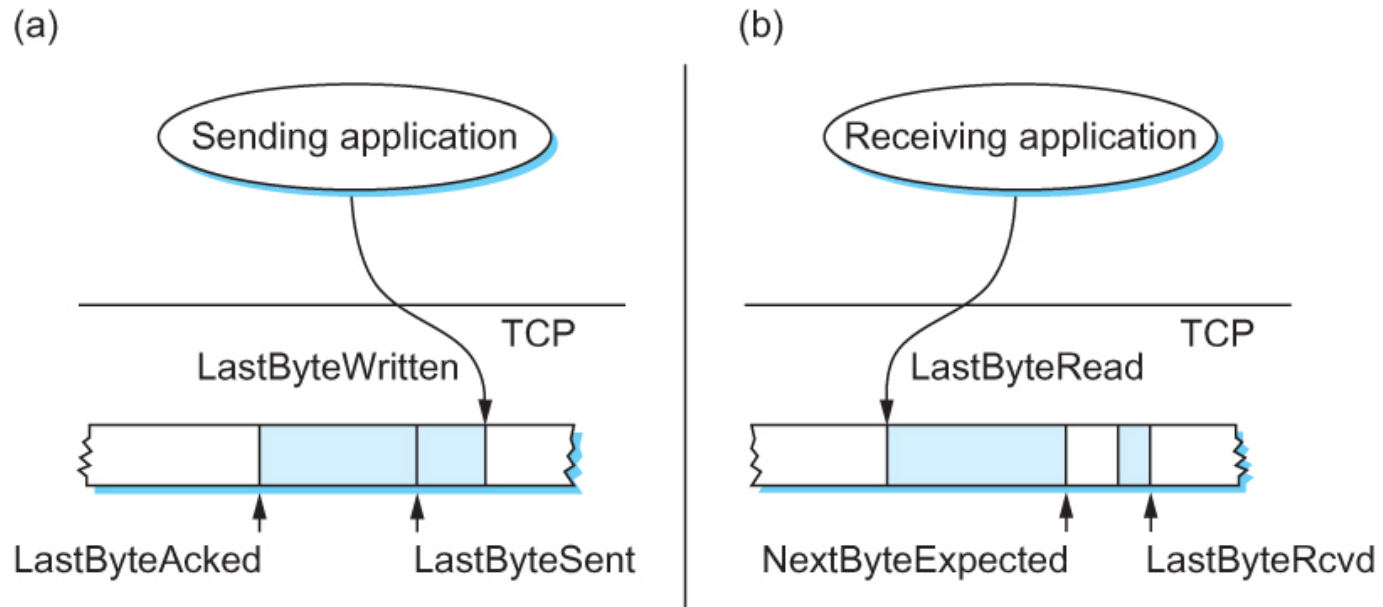
# Outline

- Simple Demultiplexer (UDP)
- Reliable Byte Stream (TCP)

# Sliding Window Revisited

- TCP's variant of the sliding window algorithm, which serves several purposes:
  - (1) it guarantees the reliable delivery of data,
  - (2) it ensures that data is delivered in order, and
  - (3) it enforces flow control between the sender and the receiver.

# Sliding Window Revisited



Relationship between TCP send buffer (a) and receive buffer (b).

# TCP Sliding Window

- Sending Side
  - $\text{LastByteAcked} \leq \text{LastByteSent}$
  - $\text{LastByteSent} \leq \text{LastByteWritten}$
- Receiving Side
  - $\text{LastByteRead} < \text{NextByteExpected}$
  - $\text{NextByteExpected} \leq \text{LastByteRcvd} + 1$

# Receiver's Advertised Window

- The big difference is the size of the sliding window size at the receiver is not fixed.
- The receiver advertises an adjustable window size (**AdvertisedWindow** field in TCP header).
- Sender is limited to having no more than **AdvertisedWindow** bytes of unACKed data at any time.

# TCP Flow Control

- Receiver throttles sender by advertising a window size no larger than the amount it can buffer.
- On TCP receiver side:  
$$\text{LastByteRcvd} - \text{LastByteRead} \leq \text{MaxRcvBuffer}$$
- to avoid buffer overflow!

# TCP Flow Control

- TCP receiver advertises:

$$\text{AdvertisedWindow} = \text{MaxRcvBuffer} - (\text{LastByteRcvd} - \text{LastByteRead})$$

- i.e., the amount of free space available in the receive buffer.



# TCP Flow Control

- TCP sender must adhere to AdvertisedWindow from the receiver such that

$\text{LastByteSent} - \text{LastByteAcked} \leq \text{AdvertisedWindow}$

- or use EffectiveWindow:

$\text{EffectiveWindow} = \text{AdvertisedWindow} - (\text{LastByteSent} - \text{LastByteAcked})$

# TCP Flow Control

- If the sending process tries to write  $y$  bytes to TCP,
- But:  
 $(\text{LastByteWritten} - \text{LastByteAked}) + y > \text{MaxSendBuffer}$

Then TCP blocks the sending process and does not allow it to generate more data.

# TCP Flow Control

- $\text{LastByteRcvd} - \text{LastByteRead} \leq \text{MaxRcvBuffer}$
- $\text{AdvertisedWindow} = \text{MaxRcvBuffer} - ((\text{NextByteExpected} - 1) - \text{LastByteRead})$
- $\text{LastByteSent} - \text{LastByteAcked} \leq \text{AdvertisedWindow}$
- $\text{EffectiveWindow} = \text{AdvertisedWindow} - (\text{LastByteSent} - \text{LastByteAcked})$
- $\text{LastByteWritten} - \text{LastByteAcked} \leq \text{MaxSendBuffer}$
- If the sending process tries to write  $y$  bytes to TCP, but  $(\text{LastByteWritten} - \text{LastByteAcked}) + y > \text{MaxSendBuffer}$  then TCP blocks the sending process and does not allow it to generate more data.

# Protecting against Wraparound

- SequenceNum: 32 bits long
- AdvertisedWindow: 16 bits long
  - TCP has satisfied the requirement of the sliding window algorithm that is the sequence number
  - space be twice as big as the window size
  - $2^{32} \gg 2 \times 2^{16}$

# Protecting against Wraparound

- Relevance of the 32-bit sequence number space
  - The sequence number used on a given connection might wraparound
  - A byte with sequence number  $x$  could be sent at one time, and then at a later time a second byte with the same sequence number  $x$  could be sent
  - Packets cannot survive in the Internet for longer than the **MSL**
  - **MSL** is set to 120 sec
  - We need to make sure that the sequence number does not wrap around within a 120-second period of time
  - Depends on how fast data can be transmitted over the Internet

# Protecting against Wraparound

Bandwidth	Time until Wraparound
T1 (1.5 Mbps)	6.4 hours
Ethernet (10 Mbps)	57 minutes
T3 (45 Mbps)	13 minutes
Fast Ethernet (100 Mbps)	6 minutes
OC-3 (155 Mbps)	4 minutes
OC-12 (622 Mbps)	55 seconds
OC-48 (2.5 Gbps)	14 seconds

Time until 32-bit sequence number space wraps around.

# Keeping the Pipe Full

- 16-bit AdvertisedWindow field must be big enough to allow the sender to keep the pipe full
- Clearly the receiver is free not to open the window as large as the AdvertisedWindow field allows
- If the receiver has enough buffer space
  - The window needs to be opened far enough to allow a full
  - delay  $\times$  bandwidth product's worth of data
  - Assuming an RTT of 100 ms

# Keeping the Pipe Full

Bandwidth	Delay $\times$ Bandwidth Product
T1 (1.5 Mbps)	18 KB
Ethernet (10 Mbps)	122 KB
T3 (45 Mbps)	549 KB
Fast Ethernet (100 Mbps)	1.2 MB
OC-3 (155 Mbps)	1.8 MB
OC-12 (622 Mbps)	7.4 MB
OC-48 (2.5 Gbps)	29.6 MB

Required window size for 100-ms RTT.



# Triggering Transmission

- How does TCP decide to transmit a segment?
  - TCP supports a byte stream abstraction
  - Application programs write bytes into streams
  - It is up to TCP to decide that it has enough bytes to send a segment

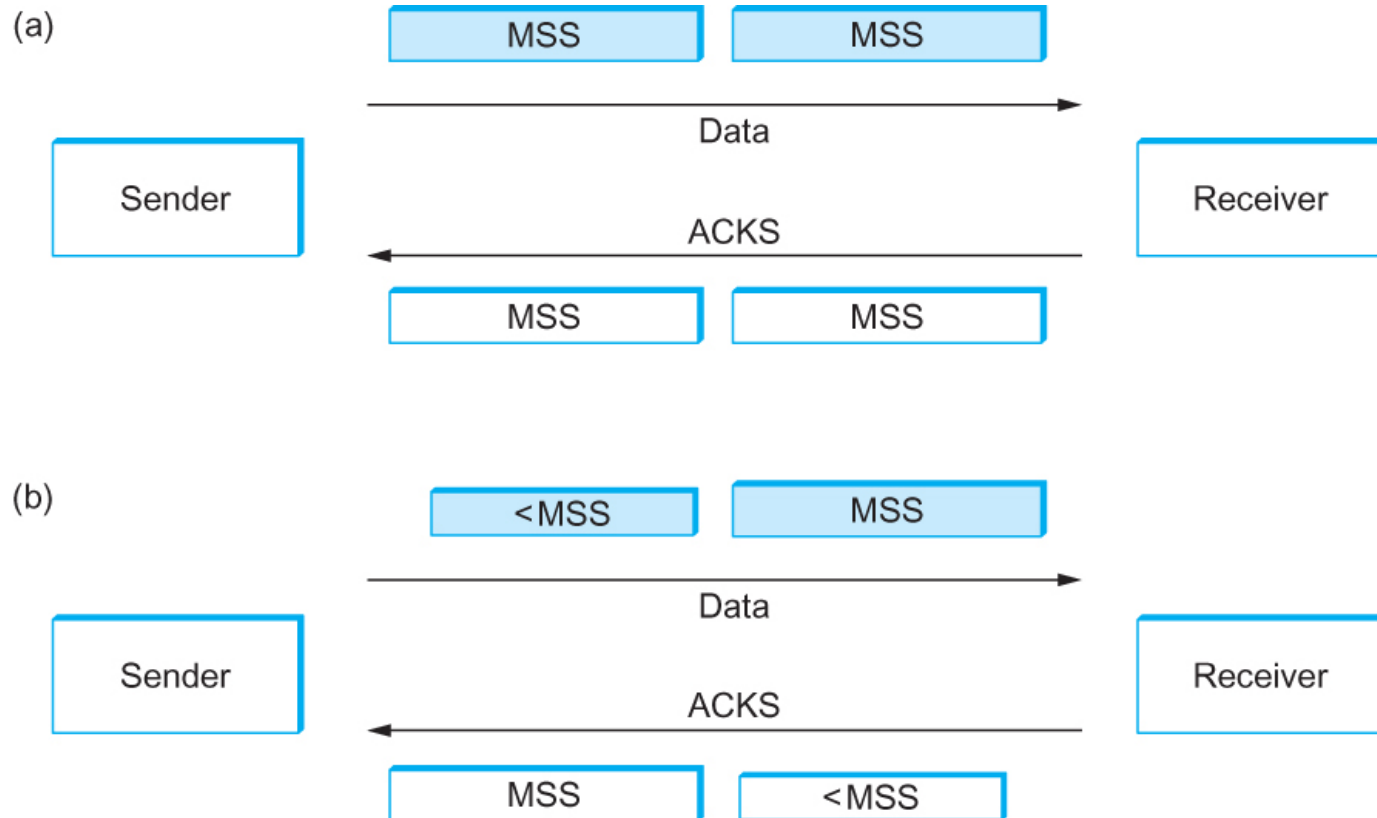
# Triggering Transmission

- What factors governs this decision
  - Ignore flow control: window is wide open, as would be the case when the connection starts
  - TCP has three mechanisms to trigger the transmission of a segment
    - 1) TCP maintains a variable MSS and sends a segment as soon as it has collected MSS bytes from the sending process
      - MSS is usually set to the size of the largest segment TCP can send without causing local IP to fragment.
      - MSS: MTU of directly connected network – (TCP header + and IP header)
    - 2) Sending process has explicitly asked TCP to send it
      - TCP supports push operation
    - 3) When a timer fires
      - Resulting segment contains as many bytes as are currently buffered for transmission

# Silly Window Syndrome

- If you think of a TCP stream as a conveyor belt with “full” containers (data segments) going in one direction and empty containers (ACKs) going in the reverse direction, then MSS-sized segments correspond to large containers and 1-byte segments correspond to very small containers.
- If the sender aggressively fills an empty container as soon as it arrives, then any small container introduced into the system remains in the system indefinitely.
- That is, it is immediately filled and emptied at each end, and never coalesced with adjacent containers to create larger containers.

# Silly Window Syndrome



Silly Window Syndrome

# Nagle's Algorithm

- If there is data to send but the window is open less than MSS, then we may want to wait some amount of time before sending the available data
- But how long?
- If we wait too long, then we hurt interactive applications like Telnet
- If we don't wait long enough, then we risk sending a bunch of tiny packets and falling into the *silly window* syndrome
  - The solution is to introduce a timer and to transmit when the timer expires

# Nagle's Algorithm

- We could use a clock-based timer, for example one that fires every 100 ms
- Nagle introduced an elegant self-clocking solution
- Key Idea
  - As long as TCP has any data in flight, the sender will eventually receive an ACK
  - This ACK can be treated like a timer firing, triggering the transmission of more data

# Nagle's Algorithm

When the application produces data to send  
    if both the available data and the window  $\geq$  MSS  
        send a full segment  
    else  
        if there is unACKed data in flight  
            buffer the new data until an ACK arrives  
        else  
            send all the new data now

# Adaptive Retransmission

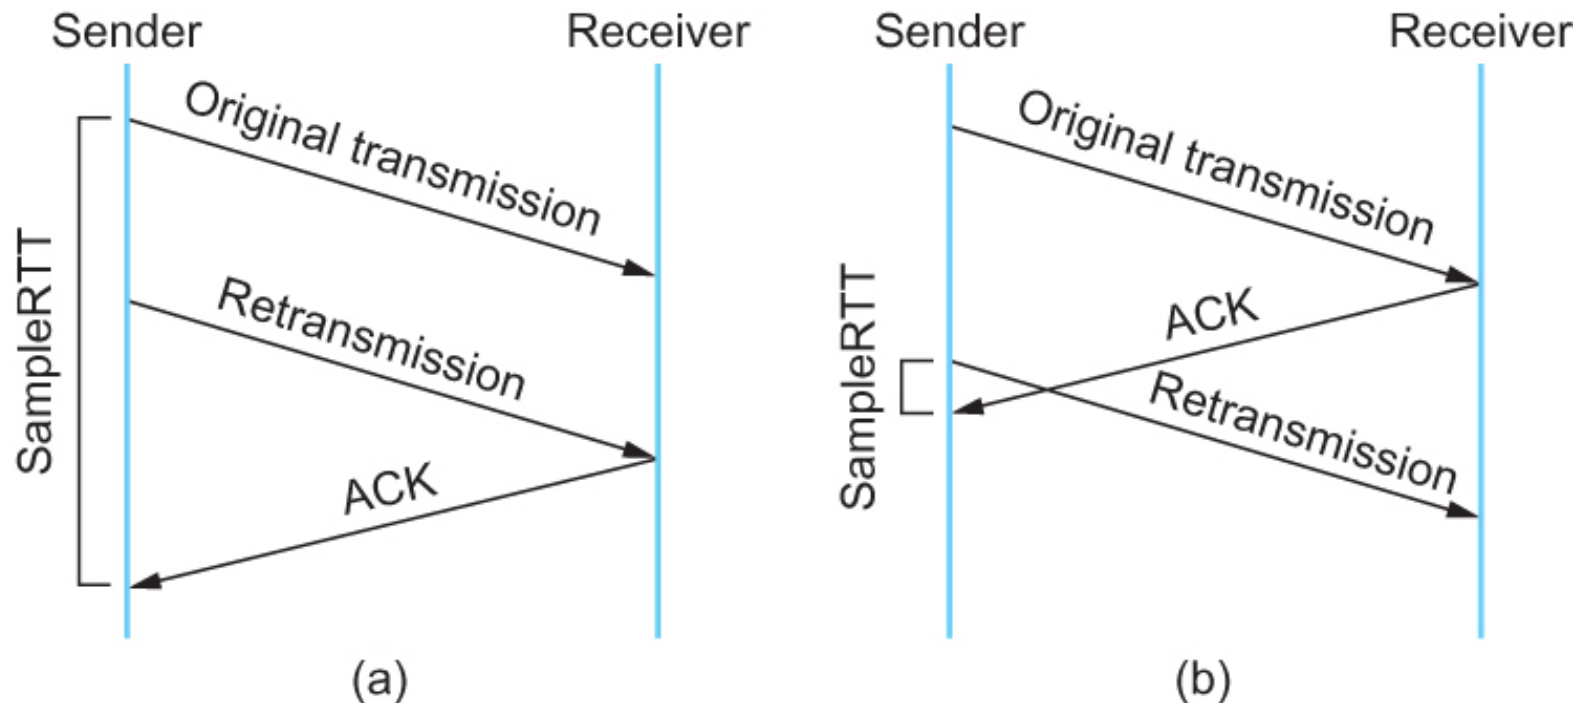
- Original Algorithm
  - Measure `sampleRTT` for each segment/ ACK pair
  - Compute weighted average of RTT
    - $\text{EstRTT} = \alpha \times \text{EstRTT} + (1 - \alpha) \times \text{SampleRTT}$ 
      - $\alpha$  between 0.8 and 0.9
  - Set timeout based on `EstRTT`
    - $\text{TimeOut} = 2 \times \text{EstRTT}$



# Original Algorithm

- Problem
  - ACK does not really acknowledge a transmission
    - It actually acknowledges the receipt of data
  - When a segment is retransmitted and then an ACK arrives at the sender
    - It is impossible to decide if this ACK should be associated with the first or the second transmission for calculating RTTs

# Karn/Partridge Algorithm



Associating the ACK with (a) original transmission versus (b) retransmission

# Karn/Partridge Algorithm

- Do not sample RTT when retransmitting
- Double timeout after each retransmission

# Karn/Partridge Algorithm

- Karn-Partridge algorithm was an improvement over the original approach, but it does not eliminate congestion
- We need to understand how timeout is related to congestion
  - If you timeout too soon, you may unnecessarily retransmit a segment which adds load to the network

# Karn/Partridge Algorithm

- Main problem with the original computation is that it does not take variance of Sample RTTs into consideration.
- If the variance among Sample RTTs is small
  - Then the Estimated RTT can be better trusted
  - There is no need to multiply this by 2 to compute the timeout

# Karn/Partridge Algorithm

- On the other hand, a large variance in the samples suggest that timeout value should not be tightly coupled to the Estimated RTT
- Jacobson/Karels proposed a new scheme for TCP retransmission

# Jacobson/Karels Algorithm

$$SRTT = \alpha SRTT + (1 - \alpha) R$$

$$RTTVAR = \beta RTTVAR + (1 - \beta) |SRTT - R|$$

$$RTO = SRTT + 4 \times RTTVAR$$

# Jacobson/Karels Algorithm

- $\text{Difference} = \text{SampleRTT} - \text{EstimatedRTT}$
- $\text{EstimatedRTT} = \text{EstimatedRTT} + (\alpha \times \text{Difference})$
- $\text{Deviation} = \text{Deviation} + (|\text{Difference}| - \text{Deviation}) \times \beta$
- $\text{TimeOut} = \mu \times \text{EstimatedRTT} + \gamma \times \text{Deviation}$ 
  - where based on experience,  $\mu$  is typically set to 1 and  $\gamma$  is set to 4. Thus, when the variance is small, TimeOut is close to EstimatedRTT; a large variance causes the deviation term to dominate the calculation.



# Summary

- We have discussed how to convert host-to-host packet delivery service to process-to-process communication channel.
- We have discussed UDP
- We have discussed TCP