

CS-AD 220 – Spring 2016

Natural Language Processing

Session 19: 7-Apr-16

Prof. Nizar Habash

NYUAD Course CS-AD 220 – Spring 2016

Natural Language Processing

Assignment #3 : POS Tagging and Parsing

Assigned Mar 31, 2016 / Due Apr 17, 2016 (11:59pm)

I. Grading & Submission

This assignment is about the development of a dependency parser and a part-of-speech (POS) tagger for English. The assignment accounts for 15% of the full grade. It consists of five exercises. **There is also a bonus exercise that can count for up to 5% of the full grade.** The additional exercise consists of a parsing competition on an unseen test set. Participation earns 2%. The first, second and third ranked systems earn additional 3%, 2% and 1%, respectively.

Assignment #3 posted on NYU Classes

START EARLY!

DEADLINE PUSHED FORWARD TO APR 17

Looking ahead

- Prof. Hend Alkhalifa
 - April 14
- Hackathon!
 - April 15-17
- Deadline Assignment #3
 - April 17
- Prof. Jan Hajic
 - April 21


Hackathon Plans

- Diplobot group
- Qusatat group
- Other groups?

Types of Syntactic Analyses

- Phrase Structure Parsing
 - aka constituency parsing
 - CKY Parser
- **Dependency Parsing**
 - **MaltParser**
- Chunking
 - Base-phrase chunking

MaltParser Example



	[ROOT]	[Happy, children, ...]
Shift	[ROOT, Happy]	[children, like, ...]
LA_{amod}	[ROOT]	[children, like, ...]
Shift	[ROOT, children]	[like, to, ...]
LA_{nsubj}	[ROOT]	[like, to, ...]
RA_{root}	[ROOT, like]	[to, play, ...]
Shift	[ROOT, like, to]	[play, with, ...]
LA_{aux}	[ROOT, like]	[play, with, ...]
RA_{xcomp}	[ROOT, like, play]	[with their, ...]

Actions	Stack	Buffer	Arcs A	Pre-conditions
Shift	σ	$w_i \beta$	A	
→	σw_i	β	A	
Reduce	σw_i	β	A	$r'(w_k, w_i) \in A$
→	σ	β	A	
Left-Arc _r	σw_i	$w_j \beta$	A	$r'(w_k, w_i) \in A, w_i \neq \text{ROOT}$
→	σ	$w_j \beta$	$A \cup \{r(w_i, w_j)\}$	e.g., children - want
Right-Arc _r	σw_i	$w_j \beta$	A	
→	$\sigma w_i w_j$	β	$A \cup \{r(w_i, w_j)\}$	e.g., want - toys

∅

∅

$\{\text{amod}(\text{children}, \text{happy})\} = A_1$

A_1

$A_1 \cup \{\text{nsubj}(\text{like}, \text{children})\} = A_2$


$A_2 \cup \{\text{root}(\text{ROOT}, \text{like})\} = A_3$

A_3

$A_3 \cup \{\text{aux}(\text{play}, \text{to})\} = A_4$

$A_4 \cup \{\text{xcomp}(\text{like}, \text{play})\} = A_5$

MaltParser Example




Actions	Stack	Buffer	Arcs A	Pre-conditions
Shift	σ	$w_i \beta$	A	
→	σw_i	β	A	
Reduce	σw_i	β	A	$r'(w_k, w_i) \in A$
→	σ	β	A	
Left-Arc _r	σw_i	$w_j \beta$	A	$r'(w_k, w_i) \in A, w_i \neq \text{ROOT}$
→	σ	$w_j \beta$	$A \cup \{r(w_i, w_j)\}$	e.g., children - want
Right-Arc _r	σw_i	$w_j \beta$	A	
→	$\sigma w_i w_j$	β	$A \cup \{r(w_i, w_j)\}$	e.g., want - toys

RA _{xcomp}	[ROOT, like, play]	[with their, ...]	$A_4 \cup \{\text{xcomp}(\text{like}, \text{play}) = A_5\}$
RA _{prep}	[ROOT, like, play, with]	[their, friends, ...]	$A_5 \cup \{\text{prep}(\text{play}, \text{with}) = A_6\}$
Shift	[ROOT, like, play, with, their]	[friends, .]	A_6
LA _{poss}	[ROOT, like, play, with]	[friends, .]	$A_6 \cup \{\text{poss}(\text{friends}, \text{their}) = A_7\}$
RA _{pobj}	[ROOT, like, play, with, friends][.]		$A_7 \cup \{\text{pobj}(\text{with}, \text{friends}) = A_8\}$
Reduce	[ROOT, like, play, with]	[.]	A_8
Reduce	[ROOT, like, play]	[.]	A_8
Reduce	[ROOT, like]	[.]	A_8
RA _{punc}	[ROOT, like, .]	[]	$A_8 \cup \{\text{punc}(\text{like}, .) = A_9\}$

Terminate as soon as the buffer is empty. Dependencies = A_9

Evaluation: Dependency Accuracy



$$\text{Acc} = \frac{\text{\# correct deps}}{\text{\# of deps}}$$

UAS =

LAS =

Gold

1	2	She	nsubj
2	0	saw	root
3	5	the	det
4	5	video	nn
5	2	lecture	dobj


Parsed

1	2	She	nsubj
2	0	saw	root
3	4	the	det
4	5	video	nsubj
5	2	lecture	ccomp

UAS = Unlabeled Attachment Score

LAS = Labeled Attachment Score

Evaluation: Dependency Accuracy



$$\text{Acc} = \frac{\text{\# correct deps}}{\text{\# of deps}}$$

$$\text{UAS} = 4 / 5 = 80\%$$

$$\text{LAS} = 2 / 5 = 40\%$$

Gold

1	2	She	nsubj
2	0	saw	root
3	5	the	det
4	5	video	nn
5	2	lecture	dobj

Parsed

1	2	She	nsubj
2	0	saw	root
3	4	the	det
4	5	video	nsubj
5	2	lecture	ccomp

UAS = Unlabeled Attachment Score

LAS = Labeled Attachment Score

Representative performance numbers

- The CoNLL-X (2006) shared task provides evaluation numbers for various dependency parsing approaches over 13 languages
 - MALT: LAS scores from 65–92%, depending greatly on language/treebank
- Here we give a few UAS numbers for English to allow some comparison to constituency parsing

Parser	UAS%
Sagae and Lavie (2006) ensemble of dependency parsers	92.7
Charniak (2000) generative, constituency	92.2
Collins (1999) generative, constituency	91.7
McDonald and Pereira (2005) – MST graph-based dependency	91.5
Yamada and Matsumoto (2003) – transition-based dependency	90.4

Projectivity

- Dependencies from a CFG tree using heads, must be **projective**
 - There must not be any crossing dependency arcs when the words are laid out in their linear order, with all arcs above the words.
- But dependency theory normally does allow non-projective structures to account for displaced constituents
 - You can't easily get the semantics of certain constructions right without these nonprojective dependencies



Handling non-projectivity

- The arc-eager algorithm we presented only builds projective dependency trees
- Possible directions to head:
 1. Just declare defeat on non-projective arcs
 2. Use a dependency formalism which only admits projective representations (a CFG doesn't represent such structures...)
 3. Use a postprocessor to a projective dependency parsing algorithm to identify and resolve non-projective links
 4. Add extra types of transitions that can model at least most non-projective structures
 5. Move to a parsing mechanism that does not use or require any constraints on projectivity (e.g., the graph-based MSTParser)

Types of Syntactic Analyses

- Phrase Structure Parsing
 - aka constituency parsing
 - CKY Parser
- Dependency Parsing
 - MaltParser
- **Chunking**
 - **Base-phrase chunking**

Phrase Chunking

- Find all non-recursive noun phrases (NPs) and verb phrases (VPs) in a sentence.
 - Base phrases
 - [NP I] [VP ate] [NP the spaghetti] [PP with] [NP meatballs].
 - [NP He] [VP reckons] [NP the current account deficit] [VP will narrow] [PP to] [NP only # 1.8 billion] [PP in] [NP September]

Phrase Chunking as Sequence Labeling

IOB Chunkers

- Tag individual words with one of 3 tags
 - B (Begin) word starts new target phrase
 - I (Inside) word is part of target phrase but not the first word
 - O (Other) word is not part of target phrase
- Sample for NP chunking
 - He reckons the current account deficit will narrow to only # 1.8 billion in September.

Begin

Inside

Other

Example

(from YAMCHA: Yet Another Multipurpose Chunk Annotator)

He	PRP	B-NP
reckons	VBZ	B-VP
the	DT	B-NP
current	JJ	I-NP
account	NN	I-NP
deficit	NN	I-NP
will	MD	B-VP
narrow	VB	I-VP
to	TO	B-PP
only	RB	B-NP
#	#	I-NP
1.8	CD	I-NP
billion	CD	I-NP
in	IN	B-PP
September	NNP	B-NP

. . O

Example

(from YAMCHA: Yet Another Multipurpose Chunk Annotator)

He	PRP	B-NP
reckons	VBZ	B-VP
the	DT	B-NP
current	JJ	I-NP
account	NN	I-NP
deficit	NN	I-NP
will	MD	B-VP
narrow	VB	I-VP
to	TO	B-PP
only	RB	B-NP
#	#	I-NP
1.8	CD	I-NP
billion	CD	I-NP
in	IN	B-PP
September	NNP	B-NP
.	.	O

*Some of the possible feature sets
that can be used to classify tags*

	COL:0	COL:1	TAG
POS:-4	He	PRP	B-NP
POS:-3	reckons	VBZ	B-VP
POS:-2	the	DT	B-NP
POS:-1	current	JJ	I-NP
POS: 0	deficit	NN	I-NP
POS:+1	will	MD	B-VP
POS:+2	narrow	VB	I-NP
POS:+3	to	TO	B-PP

	COL:0	COL:1	TAG
POS:-4	He	PRP	B-NP
POS:-3	reckons	VBZ	B-VP
POS:-2	the	DT	B-NP
POS:-1	current	JJ	I-NP
POS: 0	deficit	NN	I-NP
POS:+1	will	MD	B-VP
POS:+2	narrow	VB	I-NP
POS:+3	to	TO	B-PP

Evaluating Chunking

- Per token accuracy does not evaluate finding correct full chunks. Instead use:

$$\text{Precision} = \frac{\text{Number of correct chunks found}}{\text{Total number of chunks found}}$$

$$\text{Recall} = \frac{\text{Number of correct chunks found}}{\text{Total number of actual chunks}}$$

- Take harmonic mean to produce a single evaluation metric called F measure.

$$F = \frac{(\beta^2 + 1)PR}{\beta^2 P + R} \quad F_1 = \frac{1}{(\frac{1}{P} + \frac{1}{R})/2} = \frac{2PR}{P + R}$$

Current Chunking Results

- English NP chunking performance: $F_1 \sim 96\%$
- Typical results for finding range of chunk types (CONLL 2000 shared task: NP, VP, PP, ADV, SBAR, ADJP) is $F_1=92-94\%$

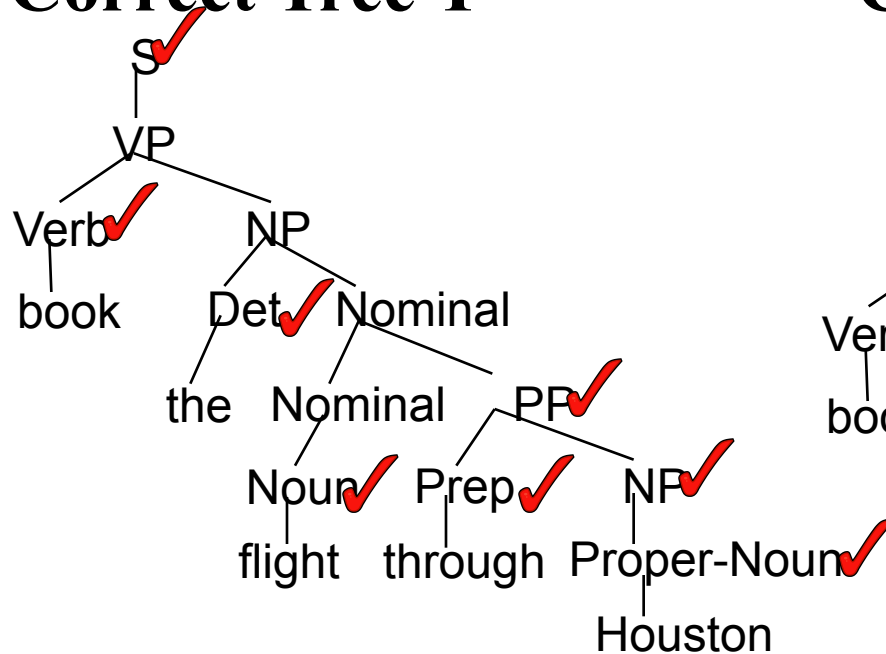
Constituency Parsing Evaluation

- PARSEVAL metrics measure the fraction of the constituents that match between the computed and human parse trees. If P is the system's parse tree and T is the human parse tree (the “gold standard”):
 - **Recall** = (# correct constituents in P) / (# constituents in T)
 - **Precision** = (# correct constituents in P) / (# constituents in P)
- **Labeled precision** and **labeled recall** require getting the non-terminal label on the constituent node correct to count as correct.
- **F_1** is the harmonic mean of precision and recall.

English labeled precision and recall number are slightly above 90%

Computing Evaluation Metrics

Correct Tree T



Constituents: 12

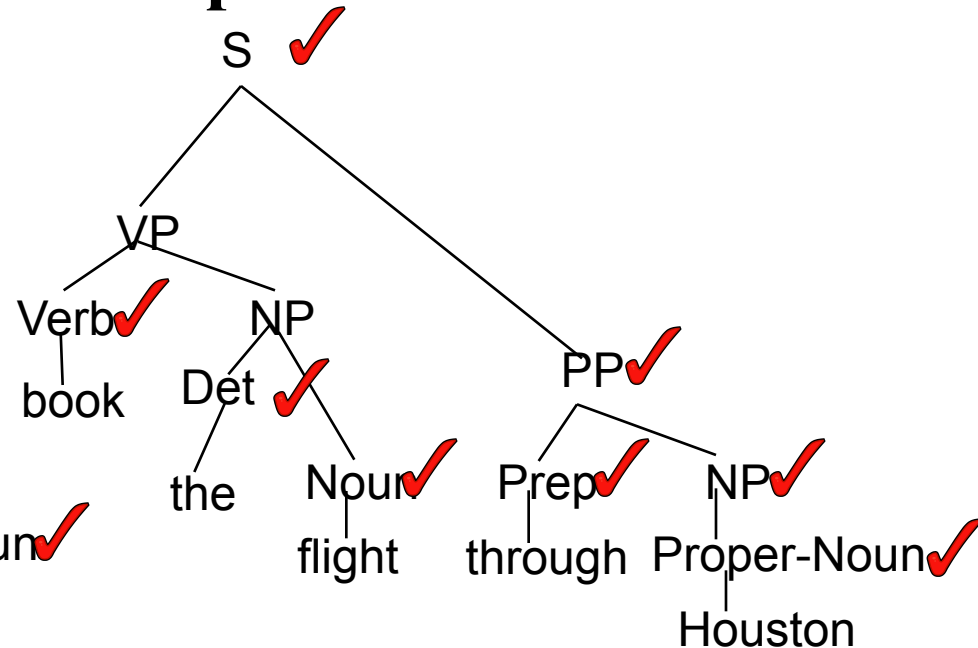
Correct Constituents: 8

Recall = $8/12 = 66.7\%$


Precision = $8/10 = 80.0\%$

$F_1 = 72.7\%$

Computed Tree P



Constituents: 10



Quiz #4

Name:_____


Create a dependency parse tree for this sentence. Use the following four relationship labels: subj, obj, mod, punc.

Retail sales drop in April cools afternoon market trading.

Quiz #4

Name:_____

Create a dependency parse tree for this sentence. Use the following four relationship labels: *subj*, *obj*, *mod*, *punc*.



Retail sales drop in April cools afternoon market trading.

Next Time

- Read J+M Chap 25 (intro up to 25.5)
- Come with questions about Assignment #3